

Terafac ML Hiring Challenge Submission

Dataset: CIFAR-10 Image Classification

Candidate: Vinay Kumar

Date: 16-01-2026

Mobile : +91 9318338110

Mail : 019vinaykumar@gmail.com

Notebook Link

<https://colab.research.google.com/drive/17uPnhD2OoIT0CiLQEGRDU74y0MVe3hOu?usp=haring>

Problem Understanding : I selected Option 1: CIFAR-10 Image Classification (10 classes, 60,000 images). Goal is to build models progressively from baseline to production-style system while maintaining clean code, reproducible experiments, and analysis.

Dataset Split

- **Train = 80%** (40,000 images from CIFAR train split)
- **Validation = 10%** (10,000 images from CIFAR train split)
- **Test = 10%** (official CIFAR-10 test split = 10,000 images)

Split method: random_split() applied on CIFAR train dataset.

Environment / Setup

Platform: Google Colab

Hardware: GPU (CUDA)

Framework: PyTorch + timm

Training style: transfer learning + ablations + ensemble + distillation

LEVEL 1 — Baseline Transfer Learning

Objective : Build a clean baseline model using transfer learning to establish benchmark accuracy.

Approach :

Used **ResNet18 pretrained** model via timm
Replaced classification head to output 10 classes
Used standard preprocessing and normalization
Saved trained weights + training plots


Model Details :


Architecture: ResNet18 (pretrained ImageNet)
Optimizer: Adam
Loss: CrossEntropyLoss
Epochs: 5
Input Size: 224×224

Result : Test Accuracy: 0.9391 (93.91%)

```
Device: cuda
Split sizes: 40000 10000 10000
Epoch 1/5 | Train Acc=0.7062 | Val Acc=0.8902
Epoch 2/5 | Train Acc=0.9067 | Val Acc=0.9227
Epoch 3/5 | Train Acc=0.9381 | Val Acc=0.9345
Epoch 4/5 | Train Acc=0.9579 | Val Acc=0.9368
Epoch 5/5 | Train Acc=0.9738 | Val Acc=0.9432

LEVEL 1 Baseline Test Accuracy: 0.9391
```





Observations :

Training curve shows stable convergence
 Confusion matrix shows most classes correctly predicted
 Minor confusion in similar classes like cat vs dog

LEVEL 2 — Intermediate Techniques (Augmentation + Tuning)



Objective : Improve baseline accuracy using augmentation + regularization and validate improvement using ablation.

Approach : Added augmentation: Horizontal flip , Rotation / crop , Color jitter

Used:

AdamW optimizer
 CosineAnnealingLR scheduler
 Label smoothing
 Conducted **ablation study**
 training without augmentation vs with augmentation

Experiments : A. without augmentation - accuracy = 95.04 , B . with augmentation - accuracy = 95.84 , **NOTE :** Augmentation increased generalization and improved test accuracy.



| LEVEL 2 Ablation Results | | |
|--------------------------|----------------------|---------------|
| | Experiment | Test Accuracy |
| 0 | Without Augmentation | 0.9504 |
| 1 | With Augmentation | 0.9584 |

LEVEL 3 — Advanced Architecture

Design + Analysis

Objective : Modify model architecture beyond baseline and add analysis such as per-class accuracy and confusion matrix.

Approach : Instead of training from scratch (which plateaued at ~85%), I applied a better approach:

- Used pretrained ResNet18 backbone
- Added SE Attention Head to improve feature learning
- Included per-class accuracy analysis

Architecture Change:

- Backbone: ResNet18 pretrained
- Added SE block to focus important channels
- Dropout head to reduce overfitting

Result : Test Accuracy: 0.9551 (95.51%)

Per-Class Performance Summary :


Strong performance across all classes.
Some weaker classes: cat (0.902), dog (0.911) due to similarity.

Insights :


SE attention helped improve class separation
Per-class breakdown shows meaningful model behavior

| Epoch | Train Acc | Val Acc |
|-------|-----------|---------|
| 1/5 | 0.8270 | 0.9256 |
| 2/5 | 0.9334 | 0.9406 |
| 3/5 | 0.9523 | 0.9487 |
| 4/5 | 0.9648 | 0.9531 |
| 5/5 | 0.9722 | 0.9543 |

LEVEL 3 Modified Architecture Test Accuracy: 0.9551



```
Per-class accuracy:  
airplane : 0.97  
automobile : 0.978  
bird : 0.959  
cat : 0.902  
deer : 0.95  
dog : 0.911  
frog : 0.98  
horse : 0.964  
ship : 0.975  
truck : 0.962
```



LEVEL 4 — Expert Technique

(Ensemble Model)

Objective : Create an ensemble pipeline to push performance using multiple trained models and a voting strategy.

Approach : Trained two different pretrained networks:

- ResNet18
- EfficientNetB0

Then combined outputs using Soft Voting Ensemble

- averaged logits: $(\text{out1} + \text{out2})/2$

Individual Model Results :

- ResNet18 Test Accuracy: **0.9556**
- EfficientNetB0 Test Accuracy: **0.9716**

Ensemble Result : Ensemble Test Accuracy: 0.9727 (97.27%)

Key Insight : Different architectures make different mistakes, so ensemble reduces error and improves accuracy.

```
Device: cuda

Model 1 Resnet18
level4_resnet18 | Epoch 1/5 | Train Acc=0.8298 | Val Acc=0.9255
level4_resnet18 | Epoch 2/5 | Train Acc=0.9356 | Val Acc=0.9440
level4_resnet18 | Epoch 3/5 | Train Acc=0.9557 | Val Acc=0.9499
level4_resnet18 | Epoch 4/5 | Train Acc=0.9678 | Val Acc=0.9533
level4_resnet18 | Epoch 5/5 | Train Acc=0.9749 | Val Acc=0.9555
Level4_resnet18 Test Accuracy: 0.9556
Saved: /content/drive/MyDrive/terafac_hackathon/models/level4_resnet18.pth

Model 2 efficientnetb0
model safetensors: 100% [██████████] 21.4M/21.4M [00:00<00:00, 28.0MB/s]
level4_effnetb0 | Epoch 1/5 | Train Acc=0.8772 | Val Acc=0.9517
level4_effnetb0 | Epoch 2/5 | Train Acc=0.9615 | Val Acc=0.9553
level4_effnetb0 | Epoch 3/5 | Train Acc=0.9809 | Val Acc=0.9664
level4_effnetb0 | Epoch 4/5 | Train Acc=0.9919 | Val Acc=0.9699
level4_effnetb0 | Epoch 5/5 | Train Acc=0.9956 | Val Acc=0.9729
level4_effnetb0 Test Accuracy: 0.9716
Saved: /content/drive/MyDrive/terafac_hackathon/models/level4_effnetb0.pth

LEVEL 4 Ensemble Test Accuracy: 0.9727

LEVEL 4 Comparison Table
  Model Test Accuracy
  0      ResNet18      0.9556
  1  EfficientNet-B0  0.9716
  2 Ensemble (Avg Logits) 0.9727
```

LEVEL 5 — Production / Research

Style System

Objective : Build a production-ready system with:

- compressed student model
- int8 quantized version
- speed evaluation
- uncertainty estimation

Approach : Step 1: Knowledge Distillation

- Teacher model: EfficientNetB0 (best performer from Level 4)
- Student model: MobileNetV3 Small (fast deployable)
- Loss: KD loss (KL divergence + CE loss)

Step 2: Quantization

- Used post-training **dynamic int8 quantization**
- Saved quantized student model

Step 3: Inference Speed Benchmark : Measured inference speed per image.

Step 4: Uncertainty Estimation : Used Monte Carlo Dropout for uncertainty demo.

Results :

- Student Test Accuracy: 0.8835
- Inference Time: 7.52 ms / image
- Quantized model saved successfully

```
Device: cuda
Loaded Teacher: /content/drive/MyDrive/terafac_hackathon/models/level4_effnetb0.pth
model.safetensors: 100% [██████████] 10.2M/10.2M [00:00<00:00, 11.5MB/s]
Epoch 1/3 | KD Loss=0.7475 | Val Acc=0.6816
Epoch 2/3 | KD Loss=0.4263 | Val Acc=0.7212
Epoch 3/3 | KD Loss=0.3545 | Val Acc=0.8782

LEVEL 5 Student (Distilled) Test Accuracy: 0.8835
Avg inference time per image (ms): 7.52
```

NOTE : Student accuracy is lower than teacher due to compression, but inference speed improved drastically and model is production-friendly.

| Uncertainty Demo: | | |
|--|------------------------|--|
| True label: 3 Pred: 3 Uncertainty score: 0.0 | | |
| LEVEL 5 Summary | | |
| | | |
| Component | Details | |
| 0 Teacher Model | EfficientNet-B0 | |
| 1 Student Model | MobileNetV3 Small (KD) | |
| 2 Quantized Student | Dynamic int8 (Linear) | |
| 3 Inference Time | 7.52 ms/img | |

Files Saved / Submission Artifacts :

Saved Models

- level1_resnet18.pth
- level2_no_aug.pth
- level2_with_aug.pth
- level3_resnet18_se.pth
- level4_resnet18.pth
- level4_effnetb0.pth
- level5_student_distilled.pth
- level5_student_quantized.pth

Stored in:

/content/drive/MyDrive/terafac_hackathon/models/

Link :

<https://drive.google.com/drive/folders/1RfB0XTytP8Yrxr-VJnQDVbKv-qMZK8wB?usp=sharing>

Requirements : torch , torchvision, timm , numpy , matplotlib , pandas , seaborn , scikit-learn

