

# Terafac ML Hiring Challenge Submission

**Dataset:** CIFAR-10 Image Classification

**Candidate:** Vinay Kumar

**Date:** 16-01-2026

**Mobile :** +91 9318338110

**Mail :** [019vinaykumar@gmail.com](mailto:019vinaykumar@gmail.com)

## Notebook Link

<https://colab.research.google.com/drive/17uPnhD2OoIT0CiLQEGRDU74y0MVe3hOu?usp=haring>

**Problem Understanding :** I selected Option 1: CIFAR-10 Image Classification (10 classes, 60,000 images). Goal is to build models progressively from baseline to production-style system while maintaining clean code, reproducible experiments, and analysis.

## Dataset Split

- **Train = 80%** (40,000 images from CIFAR train split)
- **Validation = 10%** (10,000 images from CIFAR train split)
- **Test = 10%** (official CIFAR-10 test split = 10,000 images)

Split method: random\_split() applied on CIFAR train dataset.

## Environment / Setup

Platform: Google Colab

Hardware: GPU (CUDA)

Framework: PyTorch + timm

Training style: transfer learning + ablations + ensemble + distillation

# LEVEL 1 — Baseline Transfer Learning

**Objective :** Build a clean baseline model using transfer learning to establish benchmark accuracy.

## Approach :

Used **ResNet18 pretrained** model via timm  
Replaced classification head to output 10 classes  
Used standard preprocessing and normalization  
Saved trained weights + training plots


## Model Details :


Architecture: ResNet18 (pretrained ImageNet)  
Optimizer: Adam  
Loss: CrossEntropyLoss  
Epochs: 5  
Input Size: 224×224

**Result :** Test Accuracy: 0.9391 (93.91%)

```
Device: cuda
Split sizes: 40000 10000 10000
Epoch 1/5 | Train Acc=0.7062 | Val Acc=0.8902
Epoch 2/5 | Train Acc=0.9067 | Val Acc=0.9227
Epoch 3/5 | Train Acc=0.9381 | Val Acc=0.9345
Epoch 4/5 | Train Acc=0.9579 | Val Acc=0.9368
Epoch 5/5 | Train Acc=0.9738 | Val Acc=0.9432

LEVEL 1 Baseline Test Accuracy: 0.9391
```





### Observations :

Training curve shows stable convergence  
 Confusion matrix shows most classes correctly predicted  
 Minor confusion in similar classes like cat vs dog

## LEVEL 2 — Intermediate Techniques (Augmentation + Tuning)



**Objective :** Improve baseline accuracy using augmentation + regularization and validate improvement using ablation.

**Approach :** Added augmentation: Horizontal flip , Rotation / crop , Color jitter

### Used:

AdamW optimizer  
 CosineAnnealingLR scheduler  
 Label smoothing  
 Conducted **ablation study**  
 training without augmentation vs with augmentation

**Experiments :** A. without augmentation - accuracy = 95.04 , B . with augmentation - accuracy = 95.84 , **NOTE :** Augmentation increased generalization and improved test accuracy.



LEVEL 2 Ablation Results		
	Experiment	Test Accuracy
0	Without Augmentation	0.9504
1	With Augmentation	0.9584

# LEVEL 3 — Advanced Architecture

## Design + Analysis

**Objective :** Modify model architecture beyond baseline and add analysis such as per-class accuracy and confusion matrix.

**Approach :** Instead of training from scratch (which plateaued at ~85%), I applied a better approach:

- Used pretrained ResNet18 backbone
- Added SE Attention Head to improve feature learning
- Included per-class accuracy analysis

### Architecture Change:

- Backbone: ResNet18 pretrained
- Added SE block to focus important channels
- Dropout head to reduce overfitting

**Result :** Test Accuracy: 0.9551 (95.51%)

### Per-Class Performance Summary :


Strong performance across all classes.  
Some weaker classes: cat (0.902), dog (0.911) due to similarity.

### Insights :


SE attention helped improve class separation  
Per-class breakdown shows meaningful model behavior

Epoch	Train Acc	Val Acc
1/5	0.8270	0.9256
2/5	0.9334	0.9406
3/5	0.9523	0.9487
4/5	0.9648	0.9531
5/5	0.9722	0.9543

LEVEL 3 Modified Architecture Test Accuracy: 0.9551



```
Per-class accuracy:  
airplane : 0.97  
automobile : 0.978  
bird : 0.959  
cat : 0.902  
deer : 0.95  
dog : 0.911  
frog : 0.98  
horse : 0.964  
ship : 0.975  
truck : 0.962
```



# LEVEL 4 — Expert Technique

## (Ensemble Model)

**Objective :** Create an ensemble pipeline to push performance using multiple trained models and a voting strategy.

**Approach :** Trained two different pretrained networks:

- ResNet18
- EfficientNetB0

Then combined outputs using Soft Voting Ensemble

- averaged logits:  $(\text{out1} + \text{out2})/2$

### Individual Model Results :

- ResNet18 Test Accuracy: **0.9556**
- EfficientNetB0 Test Accuracy: **0.9716**

**Ensemble Result :** Ensemble Test Accuracy: 0.9727 (97.27%)

**Key Insight :** Different architectures make different mistakes, so ensemble reduces error and improves accuracy.

```
Device: cuda

Model 1 Resnet18
level4_resnet18 | Epoch 1/5 | Train Acc=0.8298 | Val Acc=0.9255
level4_resnet18 | Epoch 2/5 | Train Acc=0.9356 | Val Acc=0.9440
level4_resnet18 | Epoch 3/5 | Train Acc=0.9557 | Val Acc=0.9499
level4_resnet18 | Epoch 4/5 | Train Acc=0.9678 | Val Acc=0.9533
level4_resnet18 | Epoch 5/5 | Train Acc=0.9749 | Val Acc=0.9555
Level4_resnet18 Test Accuracy: 0.9556
Saved: /content/drive/MyDrive/terafac_hackathon/models/level4_resnet18.pth

Model 2 efficientnetb0
model safetensors: 100% [██████████] 21.4M/21.4M [00:00<00:00, 28.0MB/s]
level4_effnetb0 | Epoch 1/5 | Train Acc=0.8772 | Val Acc=0.9517
level4_effnetb0 | Epoch 2/5 | Train Acc=0.9615 | Val Acc=0.9553
level4_effnetb0 | Epoch 3/5 | Train Acc=0.9809 | Val Acc=0.9664
level4_effnetb0 | Epoch 4/5 | Train Acc=0.9919 | Val Acc=0.9699
level4_effnetb0 | Epoch 5/5 | Train Acc=0.9956 | Val Acc=0.9729
level4_effnetb0 Test Accuracy: 0.9716
Saved: /content/drive/MyDrive/terafac_hackathon/models/level4_effnetb0.pth

LEVEL 4 Ensemble Test Accuracy: 0.9727

LEVEL 4 Comparison Table
  Model Test Accuracy
  0      ResNet18      0.9556
  1  EfficientNet-B0  0.9716
  2 Ensemble (Avg Logits)  0.9727
```

# LEVEL 5 — Production / Research

## Style System

**Objective :** Build a production-ready system with:

- compressed student model
- int8 quantized version
- speed evaluation
- uncertainty estimation

**Approach : Step 1: Knowledge Distillation**

- Teacher model: EfficientNetB0 (best performer from Level 4)
- Student model: MobileNetV3 Small (fast deployable)
- Loss: KD loss (KL divergence + CE loss)

**Step 2: Quantization**

- Used post-training **dynamic int8 quantization**
- Saved quantized student model

**Step 3: Inference Speed Benchmark :** Measured inference speed per image.

**Step 4: Uncertainty Estimation :** Used Monte Carlo Dropout for uncertainty demo.

**Results :**

- Student Test Accuracy: 0.8835
- Inference Time: 7.52 ms / image
- Quantized model saved successfully

```
Device: cuda
Loaded Teacher: /content/drive/MyDrive/terafac_hackathon/models/level4_effnetb0.pth
model.safetensors: 100% [██████████] 10.2M/10.2M [00:00<00:00, 11.5MB/s]
Epoch 1/3 | KD Loss=0.7475 | Val Acc=0.6816
Epoch 2/3 | KD Loss=0.4263 | Val Acc=0.7212
Epoch 3/3 | KD Loss=0.3545 | Val Acc=0.8782

LEVEL 5 Student (Distilled) Test Accuracy: 0.8835
Avg inference time per image (ms): 7.52
```

**NOTE :** Student accuracy is lower than teacher due to compression, but inference speed improved drastically and model is production-friendly.

Uncertainty Demo:

True label: 3 | Pred: 3 | Uncertainty score: 0.0

#### LEVEL 5 Summary

Component	Details
0 Teacher Model	EfficientNet-B0
1 Student Model	MobileNetV3 Small (KD)
2 Quantized Student	Dynamic int8 (Linear)
3 Inference Time	7.52 ms/img

